

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

C:\Users\hp\anaconda3\lib\site-packages\scipy\_init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.4
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
In [2]: df = pd.read_excel("C:/Users/hp/Downloads/WHO-UCN-GTB-PCI-2021.7-eng.xlsx", sheet_name = 'Mutation_catalogue')
```

```
In [3]: df.head()
```

Out[3]:

	Drug	Tier	Common Variant	Genome position	algorithm pass	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	...	RIF CC guide 2021	GenoTy MTBDRpl V;
0	AMI	1.0	rrs_a1401g	1473246.0	1.0	918.0	966.0	50.0	15640.0	939.0	...	NaN	N;
1	AMI	1.0	eis_c-14t	2715346.0	1.0	28.0	77.0	51.0	7325.0	32.0	...	NaN	N;
2	AMI	1.0	rrs_g1484t	1473329.0	1.0	5.0	7.0	2.0	15688.0	6.0	...	NaN	N;
3	AMI	1.0	rrs_c1402t	1473247.0	1.0	4.0	13.0	10.0	15680.0	5.0	...	NaN	N;
4	AMI	2.0	whiB6_A77V	NaN	1.0	3.0	100.0	97.0	15593.0	3.0	...	NaN	N;

5 rows × 52 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17396 entries, 0 to 17395
Data columns (total 52 columns):
```

#	Column	Non-Null Count	Dtype
0	Drug	17396 non-null	object
1	Tier	17348 non-null	float64
2	Common Variant	17396 non-null	object
3	Genome position	6486 non-null	float64
4	algorithm pass	17388 non-null	float64
5	Present_SOLO_R	17388 non-null	float64
6	Present_SOLO_SR	17388 non-null	float64
7	Present_S	17388 non-null	float64
8	Absent_S	17388 non-null	float64
9	Present_R	17388 non-null	float64
10	Absent_R	17388 non-null	float64
11	PPV	17388 non-null	float64
12	PPV_lb	17388 non-null	float64
13	PPV_ub	17388 non-null	float64
14	PPV SOLO	16074 non-null	float64
15	PPV SOLO_lb	16074 non-null	float64
16	PPV SOLO_ub	16074 non-null	float64
17	Sensitivity	17388 non-null	float64
18	Sensitivity_lb	17388 non-null	float64
19	Sensitivity_ub	17388 non-null	float64
20	Specificity	17388 non-null	float64
21	Specificity_lb	17388 non-null	float64
22	Specificity_ub	17388 non-null	float64
23	LR+	17388 non-null	float64
24	LR+_lb	17388 non-null	float64
25	LR+_ub	17388 non-null	float64
26	LR-	17385 non-null	float64
27	LR-_lb	17388 non-null	float64
28	LR-_ub	17388 non-null	float64
29	OR	17385 non-null	float64
30	OR_lb	17388 non-null	float64
31	OR_ub	17388 non-null	float64
32	OR SOLO	7696 non-null	float64
33	OR SOLO_lb	7696 non-null	float64
34	OR SOLO_ub	7696 non-null	float64
35	OR SOLO_FE-sig	17388 non-null	float64
36	Neutral masked	17388 non-null	float64
37	INITIAL CONFIDENCE GRADING	17388 non-null	object
38	DATASET(S)	17388 non-null	object
39	Miotto et al. (PMID 29284687)	311 non-null	object
40	NGS Guide 2018	256 non-null	object
41	Level of resistance to INH or MXF	131 non-null	object
42	RIF CC guide 2021	135 non-null	object
43	Hain GenoType MTBDRplus V2.0	151 non-null	object
44	Nipro Genoscholar NTM+MDRTB II	176 non-null	object
45	Cepheid Xpert MTB/RIF	135 non-null	object
46	Cepheid Xpert MTB/RIF Ultra	139 non-null	object
47	Hain GenoType MTBDRsl V2.0	76 non-null	object
48	Cepheid Xpert MTB/XDR	199 non-null	object
49	Nipro Genoscholar PZA-TB II	503 non-null	object
50	Additional grading criteria	936 non-null	object
51	FINAL CONFIDENCE GRADING	17396 non-null	object

dtypes: float64(35), object(17)
memory usage: 6.9+ MB

```
In [5]: df.isnull().sum()#Checking for missing values
```

```
Out[5]: Drug 0
Tier 48
Common Variant 0
Genome position 10910
algorithm pass 8
Present_SOLO_R 8
Present_SOLO_SR 8
Present_S 8
Absent_S 8
Present_R 8
Absent_R 8
PPV 8
PPV_lb 8
PPV_ub 8
PPV | SOLO 1322
PPV | SOLO_lb 1322
PPV | SOLO_ub 1322
Sensitivity 8
Sensitivity_lb 8
Sensitivity_ub 8
Specificity 8
Specificity_lb 8
Specificity_ub 8
LR+ 8
LR+_lb 8
LR+_ub 8
LR- 11
LR-_lb 8
LR-_ub 8
OR 11
OR_lb 8
OR_ub 8
OR SOLO 9700
OR SOLO_lb 9700
OR SOLO_ub 9700
OR SOLO_FE-sig 8
Neutral masked 8
INITIAL CONFIDENCE GRADING 8
DATASET(S) 8
Miotto et al. (PMID 29284687) 17085
NGS Guide 2018 17140
Level of resistance to INH or MXF 17265
RIF CC guide 2021 17261
Hain GenoType MTBDRplus V2.0 17245
Nipro Genoscholar NTM+MDRTB II 17220
Cepheid Xpert MTB/RIF 17261
Cepheid Xpert MTB/RIF Ultra 17257
Hain GenoType MTBDRsl V2.0 17320
Cepheid Xpert MTB/XDR 17197
Nipro Genoscholar PZA-TB II 16893
Additional grading criteria 16460
FINAL CONFIDENCE GRADING 0
dtype: int64
```

```
In [6]: df.columns #Checking for the number of features
```

```
Out[6]: Index(['Drug', 'Tier', 'Common Variant', 'Genome position', 'algorithm pass',
      'Present_SOLO_R', 'Present_SOLO_SR', 'Present_S', 'Absent_S',
      'Present_R', 'Absent_R', 'PPV', 'PPV_lb', 'PPV_ub', 'PPV | SOLO',
      'PPV | SOLO_lb', 'PPV | SOLO_ub', 'Sensitivity', 'Sensitivity_lb',
      'Sensitivity_ub', 'Specificity', 'Specificity_lb', 'Specificity_ub',
      'LR+', 'LR+_lb', 'LR+_ub', 'LR-', 'LR-_lb', 'LR-_ub', 'OR', 'OR_lb',
      'OR_ub', 'OR SOLO', 'OR SOLO_lb', 'OR SOLO_ub', 'OR SOLO_FE-sig',
      'Neutral masked', 'INITIAL CONFIDENCE GRADING', 'DATASET(S)',
      'Miotto et al. (PMID 29284687)', 'NGS Guide 2018',
      'Level of resistance to INH or MXF', 'RIF CC guide 2021',
      'Hain GenoType MTBDRplus V2.0', 'Nipro Genoscholar NTM+MDRTB II',
      'Cepheid Xpert MTB/RIF', 'Cepheid Xpert MTB/RIF Ultra',
      'Hain GenoType MTBDRsl V2.0', 'Cepheid Xpert MTB/XDR',
      'Nipro Genoscholar PZA-TB II', 'Additional grading criteria',
      'FINAL CONFIDENCE GRADING'],
      dtype='object')
```

```
In [7]: drop_columns = ['Genome position', 'algorithm pass', 'PPV', 'PPV_lb', 'PPV_ub', 'PPV | SOLO',
      'PPV | SOLO_lb', 'PPV | SOLO_ub', 'Sensitivity', 'Sensitivity_lb', 'Sensitivity_ub',
      'Specificity', 'Specificity_lb', 'Specificity_ub', 'LR+', 'LR+_lb', 'LR+_ub', 'LR-', 'LR-_lb',
      'LR-_ub', 'OR', 'OR_lb', 'OR_ub', 'OR SOLO', 'OR SOLO_lb', 'OR SOLO_ub', 'OR SOLO_FE-sig',
      'DATASET(S)', 'Miotto et al. (PMID 29284687)', 'NGS Guide 2018', 'Level of resistance to INH or
      RIF CC guide 2021', 'Hain GenoType MTBDRplus V2.0', 'Nipro Genoscholar NTM+MDRTB II',
      'Cepheid Xpert MTB/RIF', 'Cepheid Xpert MTB/RIF Ultra', 'Hain GenoType MTBDRsl V2.0',
      'Cepheid Xpert MTB/XDR', 'Nipro Genoscholar PZA-TB II', 'Additional grading criteria']

# Dropping the specified columns from the DataFrame
df.drop(columns=drop_columns, inplace=True)
```

```
In [8]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17396 entries, 0 to 17395
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Drug                                  17396 non-null  object
1   Tier                                  17348 non-null  float64
2   Common Variant                       17396 non-null  object
3   Present_SOLO_R                       17388 non-null  float64
4   Present_SOLO_SR                      17388 non-null  float64
5   Present_S                             17388 non-null  float64
6   Absent_S                             17388 non-null  float64
7   Present_R                             17388 non-null  float64
8   Absent_R                             17388 non-null  float64
9   Neutral masked                       17388 non-null  float64
10  INITIAL CONFIDENCE GRADING           17388 non-null  object
11  FINAL CONFIDENCE GRADING             17396 non-null  object
dtypes: float64(8), object(4)
memory usage: 1.6+ MB

```

```
In [9]: df.isnull().sum()
```

```

Out[9]: Drug                                0
Tier                                48
Common Variant                        0
Present_SOLO_R                       8
Present_SOLO_SR                      8
Present_S                             8
Absent_S                             8
Present_R                             8
Absent_R                             8
Neutral masked                       8
INITIAL CONFIDENCE GRADING           8
FINAL CONFIDENCE GRADING             0
dtype: int64

```

```
In [10]: df.dtypes
```

```

Out[10]: Drug                                object
Tier                                float64
Common Variant                        object
Present_SOLO_R                       float64
Present_SOLO_SR                      float64
Present_S                             float64
Absent_S                             float64
Present_R                             float64
Absent_R                             float64
Neutral masked                       float64
INITIAL CONFIDENCE GRADING           object
FINAL CONFIDENCE GRADING             object
dtype: object

```

```

In [11]: #Handle missing values in categorical columns by imputing with mode
categorical_cols = ['Drug', 'Common Variant', 'Neutral masked', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']
df[categorical_cols] = df[categorical_cols].fillna(df[categorical_cols].mode().iloc[0])

df.head()

```

```

Out[11]:
   Drug Tier Common Variant Present_SOLO_R Present_SOLO_SR Present_S Absent_S Present_R Absent_R Neutral masked INITIAL CONFIDENCE GRADING FINAL CONFIDENCE GRADING
0  AMI   1.0  rrs_a1401g         918.0         966.0        50.0   15640.0      939.0     349.0         0.0      Assoc w R      1) Assc
1  AMI   1.0   eis_c-14t          28.0          77.0        51.0    7325.0       32.0     632.0         0.0      Assoc w R      1) Assc
2  AMI   1.0  rrs_g1484t           5.0           7.0         2.0   15688.0        6.0    1282.0         0.0      Assoc w R      2) Assc
3  AMI   1.0  rrs_c1402t           4.0          13.0        10.0   15680.0        5.0    1283.0         0.0  Uncertain significance  2) Assc
4  AMI   2.0 whiB6_A77V            3.0         100.0       97.0   15593.0        3.0     141.0         0.0  Uncertain significance  3) Ur sign

```

```
In [12]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17396 entries, 0 to 17395
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Drug                                  17396 non-null  object
1   Tier                                 17348 non-null  float64
2   Common Variant                       17396 non-null  object
3   Present_SOLO_R                       17388 non-null  float64
4   Present_SOLO_SR                      17388 non-null  float64
5   Present_S                             17388 non-null  float64
6   Absent_S                             17388 non-null  float64
7   Present_R                             17388 non-null  float64
8   Absent_R                             17388 non-null  float64
9   Neutral masked                       17396 non-null  float64
10  INITIAL CONFIDENCE GRADING           17396 non-null  object
11  FINAL CONFIDENCE GRADING             17396 non-null  object
dtypes: float64(8), object(4)
memory usage: 1.6+ MB

```

```

In [13]: unique_values = ['Drug', 'Common Variant', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']

```

```

# Calculate the number of unique values for the specified columns
unique_count = df[unique_values].nunique()

unique_count

```

```

Out[13]: Drug                15
Common Variant          13449
INITIAL CONFIDENCE GRADING      6
FINAL CONFIDENCE GRADING      6
dtype: int64

```

```


In [14]: numeric_columns = ['Tier', 'Present_SOLO_R', 'Present_SOLO_SR', 'Present_S', 'Absent_S',
                             'Present_R', 'Absent_R']

```

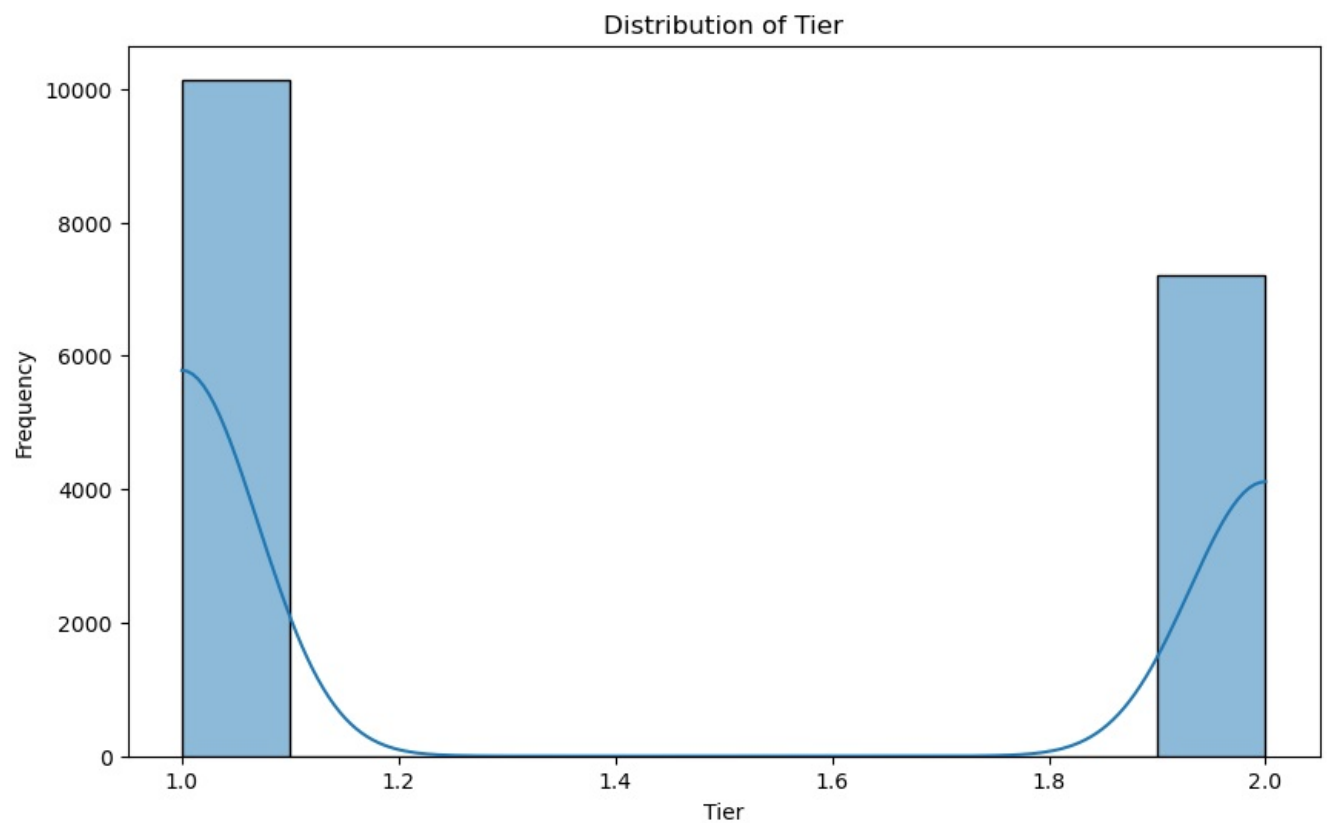
```

In [15]: sns.pairplot(df[numeric_columns])
plt.show()

```




```
In [16]: plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='Tier', bins=10, kde=True)
plt.xlabel('Tier')
plt.ylabel('Frequency')
plt.title('Distribution of Tier')
plt.show()
```



```
In [17]: plt.figure(figsize=(10, 6))
ax = sns.countplot(data=df, x='Drug', palette='bright')

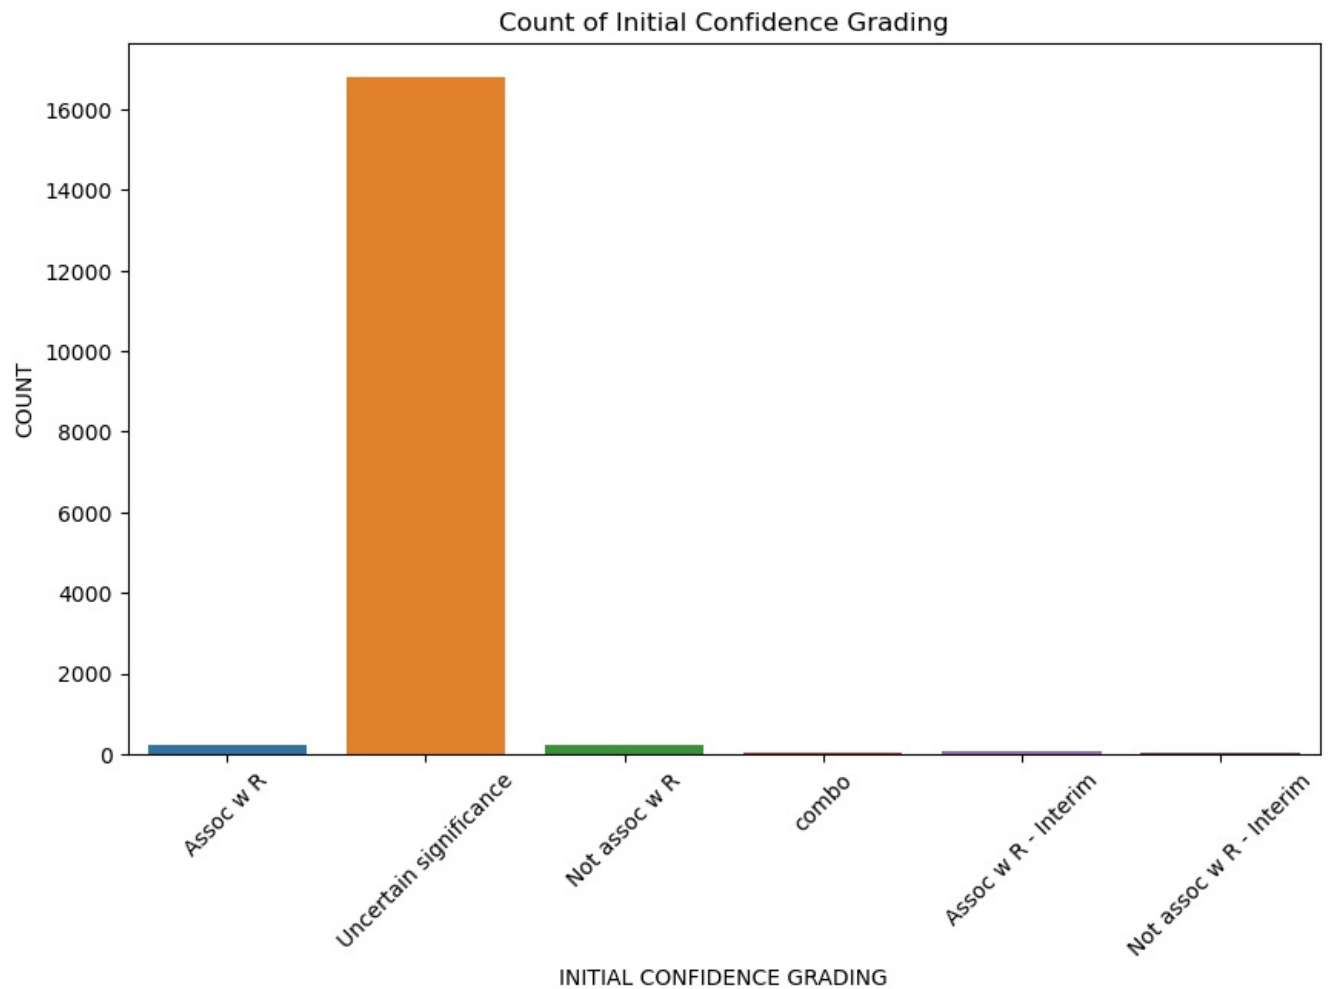
# Add text annotations to the bars
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 5),
                textcoords='offset points')

plt.xticks(rotation=45)
plt.xlabel('Drug')
plt.ylabel('Count')
plt.title('Count of Each Drug')
plt.show()
```

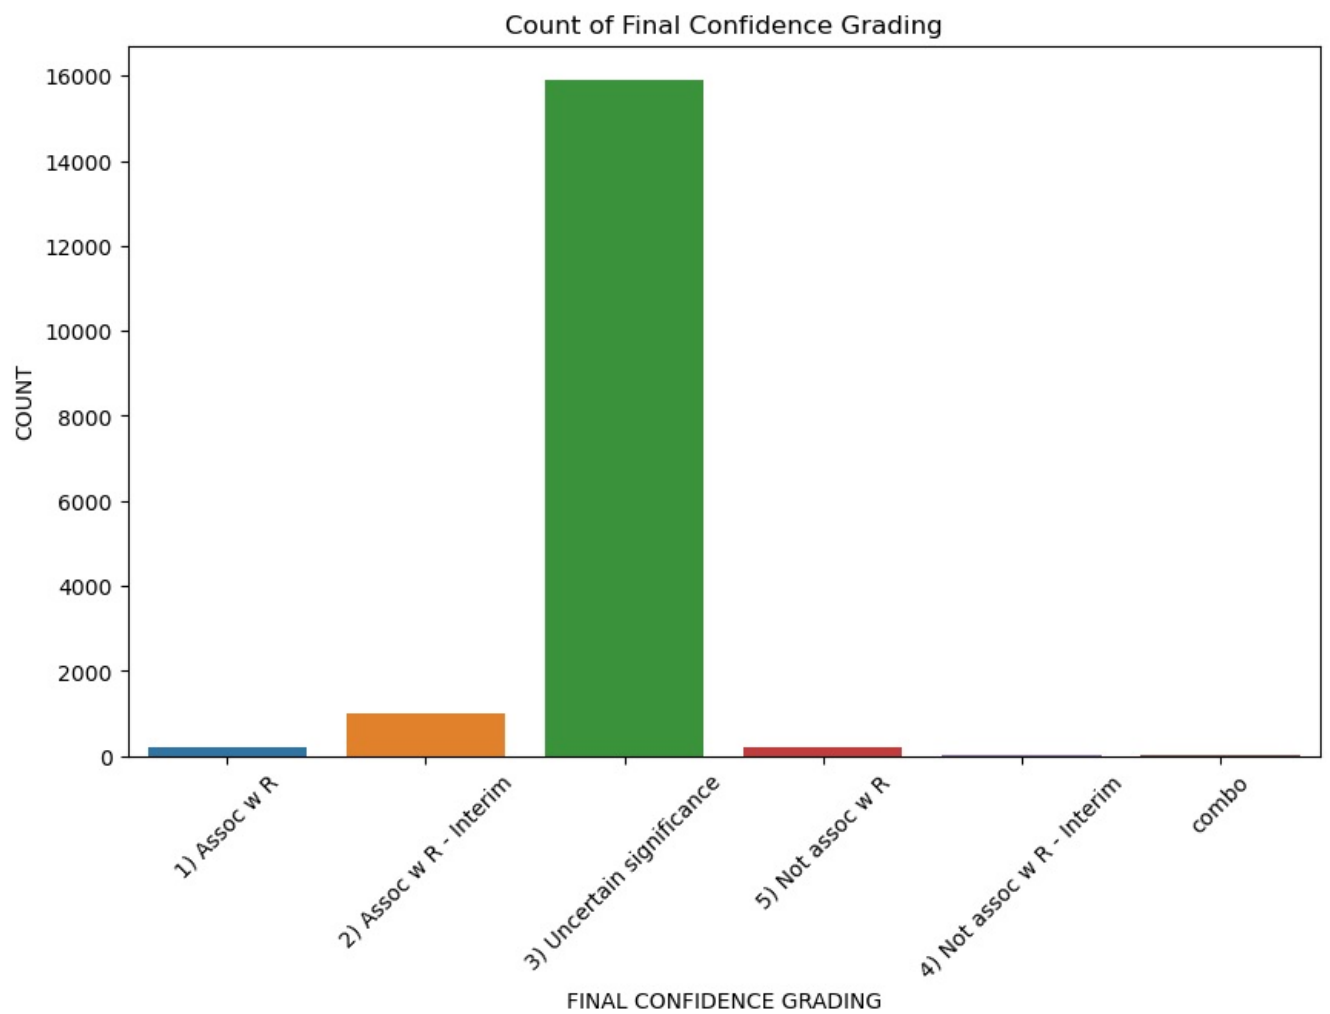


```
In [18]: plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='INITIAL CONFIDENCE GRADING')
```

```
plt.xticks(rotation=45)
plt.xlabel('INITIAL CONFIDENCE GRADING')
plt.ylabel('COUNT')
plt.title('Count of Initial Confidence Grading')
plt.show()
```

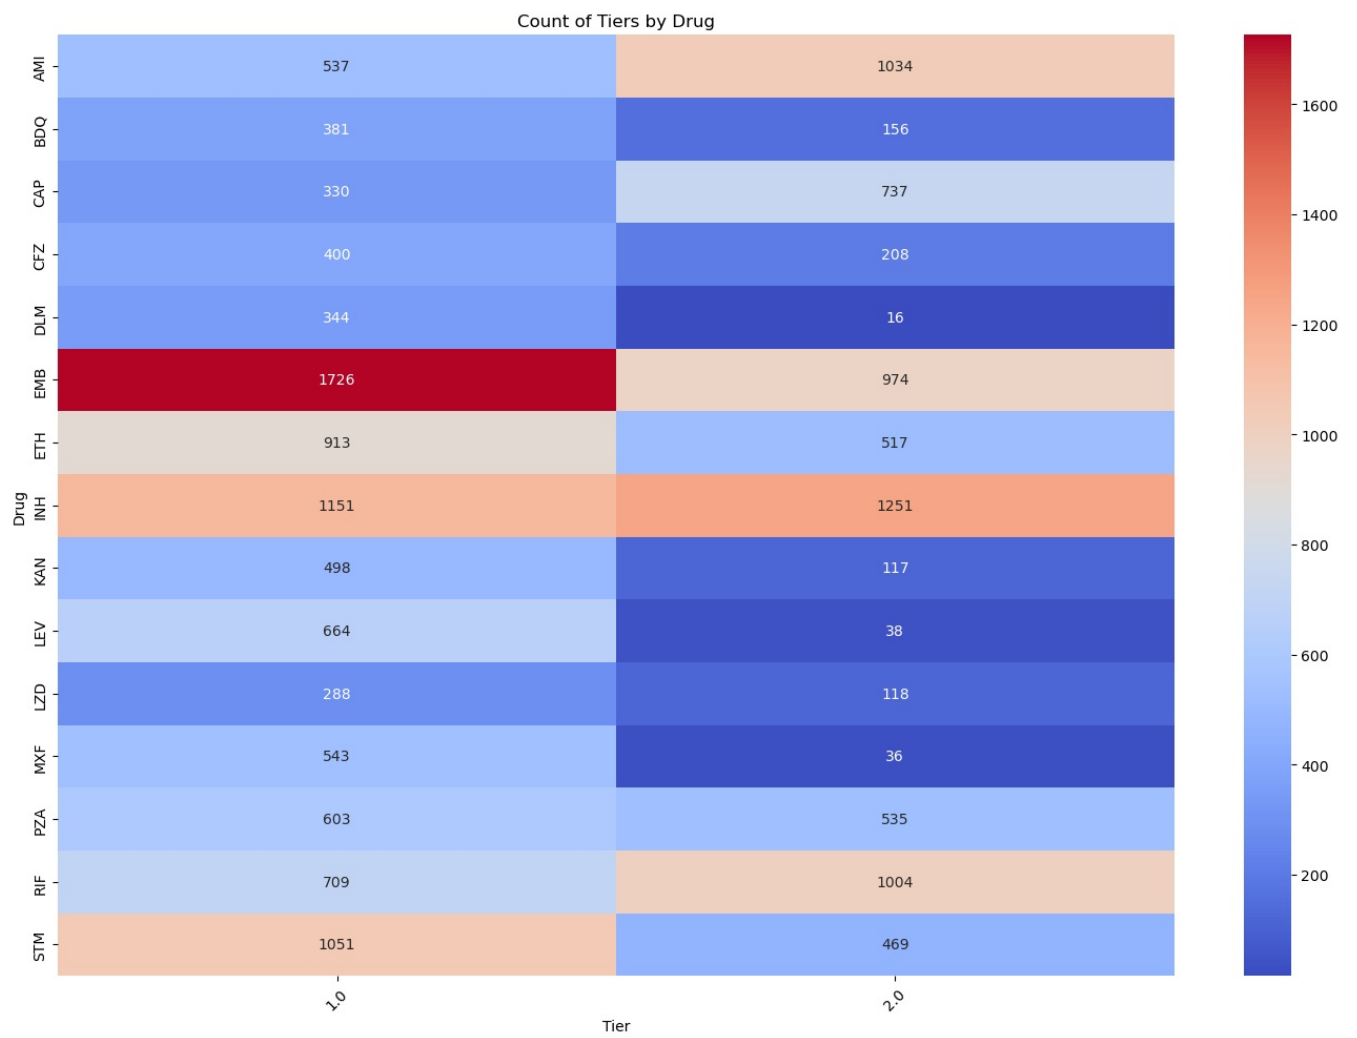


```
In [19]: plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='FINAL CONFIDENCE GRADING')
plt.xticks(rotation=45)
plt.xlabel('FINAL CONFIDENCE GRADING')
plt.ylabel('COUNT')
plt.title('Count of Final Confidence Grading')
plt.show()
```


```
In [20]: pivot_table = df.pivot_table(index='Drug', columns='Tier', aggfunc='size', fill_value=0)

# Plot heatmap
plt.figure(figsize=(14, 10))
sns.heatmap(pivot_table, cmap='coolwarm', annot=True, fmt='d')
plt.title('Count of Tiers by Drug')
plt.xlabel('Tier')
plt.ylabel('Drug')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

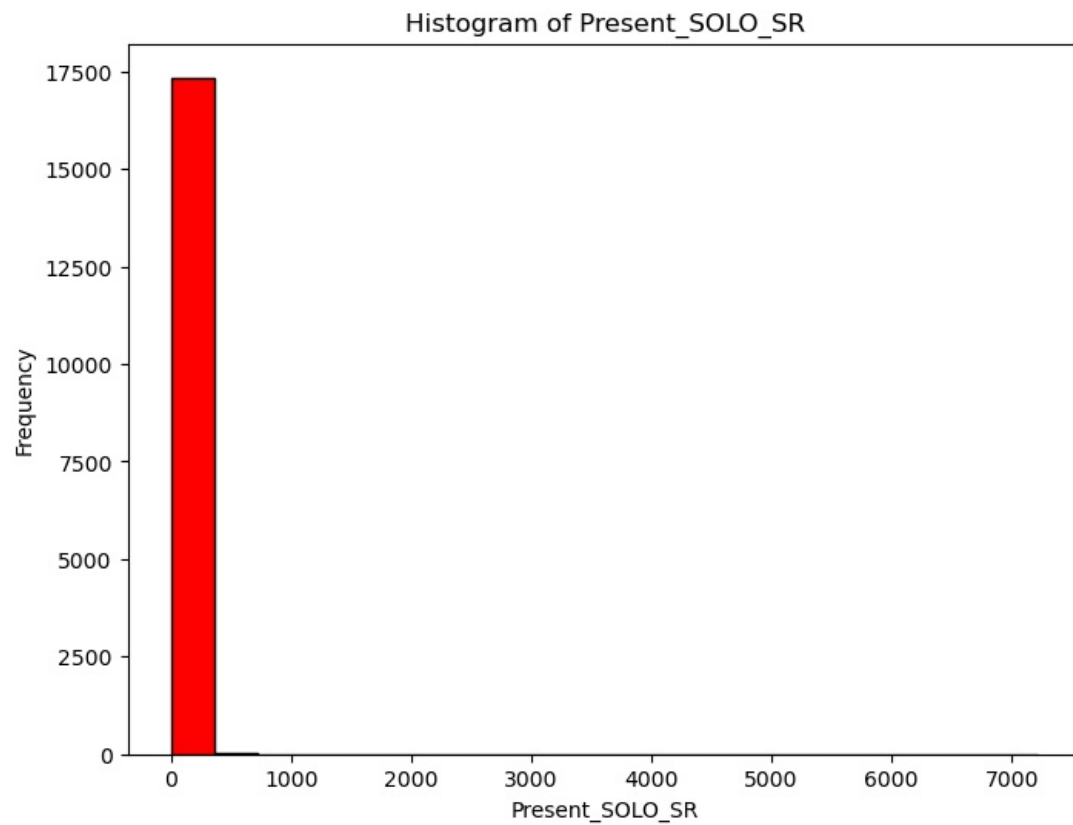


```
In [21]: from sklearn.preprocessing import LabelEncoder
#categorical_cols = df.select_dtypes(include='object').columns
```


```
In [22]: plt.figure(figsize=(8, 6))
plt.hist(df['Present_SOLO_R'], bins=20, color='green', edgecolor='black')
plt.title('Histogram of Present_SOLO_R')
plt.xlabel('Present_SOLO_R')
plt.ylabel('Frequency')
plt.grid(False)
plt.show()
```



```
In [23]: plt.figure(figsize=(8, 6))
plt.hist(df['Present_SOLO_SR'], bins=20, color='red', edgecolor='black')
plt.title('Histogram of Present_SOLO_SR')
plt.xlabel('Present_SOLO_SR')
plt.ylabel('Frequency')
plt.grid(False)
plt.show()
```

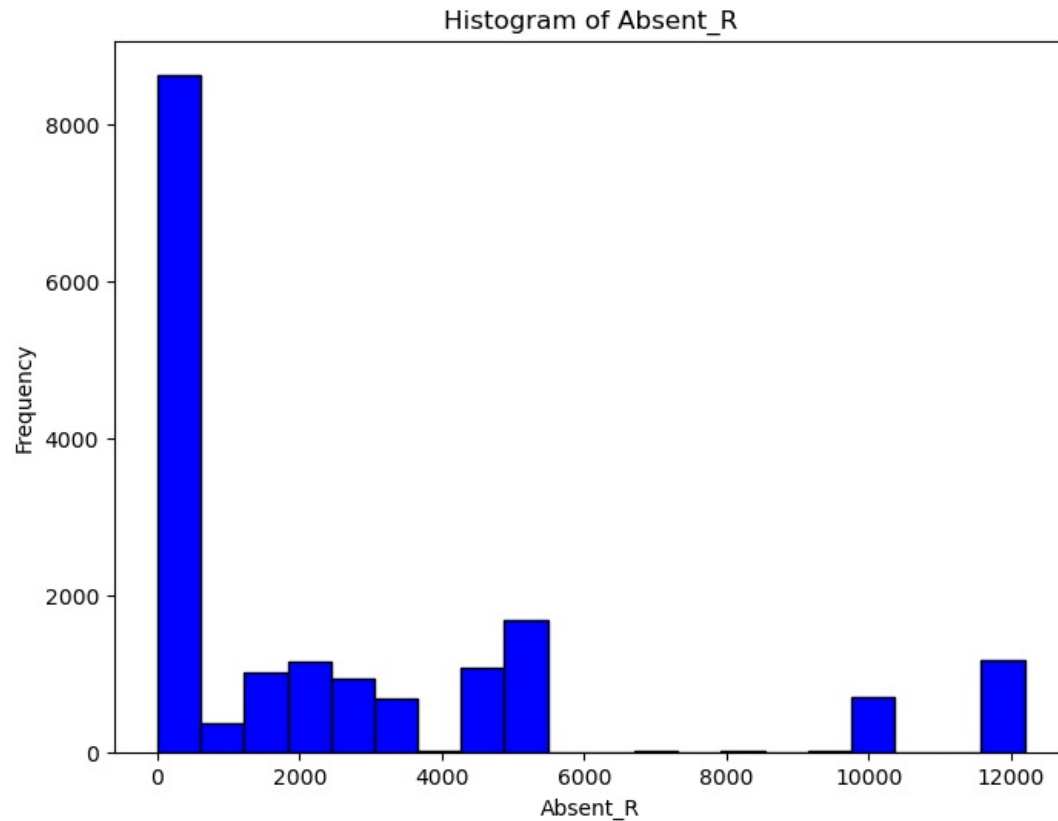


```
In [24]: plt.figure(figsize=(8, 6))
plt.hist(df['Present_S'], bins=20, color='yellow', edgecolor='black')
plt.title('Histogram of Present_S')
plt.xlabel('Present_S')
plt.ylabel('Frequency')
plt.grid(False)
plt.show()
```



```
In [25]: plt.figure(figsize=(8, 6))
plt.hist(df['Absent_R'], bins=20, color='blue', edgecolor='black')
```

```
plt.title('Histogram of Absent_R')
plt.xlabel('Absent_R')
plt.ylabel('Frequency')
plt.grid(False)
plt.show()
```



```
In [26]: # unique values and their counts in the 'Absent_R' feature
unique_values_counts = df['Absent_R'].value_counts()

unique_values_counts
```

```
Out[26]: 4900.0    1405
480.0      1173
144.0       991
321.0       964
229.0       932
...
12147.0      1
12145.0      1
12142.0      1
12139.0      1
2543.0       1
Name: Absent_R, Length: 535, dtype: int64
```

```
In [ ]: import numpy as np
        Replace infinity values with a very large number
        df.replace([np.inf, -np.inf], np.finfo(np.float64).max, inplace=True)
```

```
In [ ]: #from sklearn.preprocessing import OneHotEncoder

        # Select categorical columns for one-hot encoding
        #categorical_cols = ['Drug', 'Common Variant', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']

        # Perform one-hot encoding
        #encoded_df = pd.get_dummies(df, columns=categorical_cols)

        # Display the encoded DataFrame
        #print(encoded_df)
```

```
In [ ]: from sklearn.preprocessing import LabelEncoder

        # LabelEncoder
        encoder = LabelEncoder()

        #Encode categorical columns
        categorical_cols = ['Drug', 'Common Variant', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']
        df[categorical_cols] = df[categorical_cols].apply(encoder.fit_transform)
```

```
In [ ]: #import pandas as pd
        #from sklearn.preprocessing import OneHotEncoder

        # Assuming X_encoded contains the one-hot encoded features and y contains the target variable
```

```
# Concatenate X_encoded with the target variable y
#encoded_df = pd.concat([X_encoded, y], axis=1)

# Now encoded_df contains the one-hot encoded features along with the target variable
# You can use encoded_df for further analysis, modeling, etc.
```

In [57]: `df.head()`

Out[57]:

	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R	Neutral masked	INITIAL CONFIDENCE GRADING	FINAL CONFIDENCE GRADING
0	0	0.0	11986	0.128571	0.133776	0.005898	0.606014	0.099019	0.028618	0.0	0	
1	0	0.0	2333	0.003922	0.010663	0.006016	0.283827	0.003374	0.051825	0.0	0	
2	0	0.0	12271	0.000700	0.000969	0.000236	0.607874	0.000633	0.105125	0.0	0	
3	0	0.0	12111	0.000560	0.001800	0.001180	0.607564	0.000527	0.105207	0.0	4	
4	0	1.0	12971	0.000420	0.013848	0.011443	0.604192	0.000316	0.011562	0.0	4	

In [58]: `df.tail()`

Out[58]:

	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R	Neutral masked	INITIAL CONFIDENCE GRADING	FINAL CONFIDENCE GRADING
17391	14	0.0	13290	0.00028	0.036837	0.033620	0.240081	0.002109	0.208446	1.0	2	
17392	14	0.0	12149	0.00000	0.000000	0.042350	0.237213	0.002004	0.208528	1.0	2	
17393	14	1.0	11603	0.00000	0.000000	0.005898	0.249186	0.000000	0.006970	1.0	2	
17394	14	1.0	12784	0.00000	0.000000	0.007432	0.248683	0.000000	0.006970	1.0	2	
17395	14	NaN	11566	0.00000	0.000000	0.000000	0.362252	0.000844	0.379418	0.0	5	

In [59]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17396 entries, 0 to 17395
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Drug                                  17396 non-null  int32
1   Tier                                  17348 non-null  float64
2   Common Variant                       17396 non-null  int32
3   Present_SOLO_R                       17388 non-null  float64
4   Present_SOLO_SR                      17388 non-null  float64
5   Present_S                            17388 non-null  float64
6   Absent_S                             17388 non-null  float64
7   Present_R                            17388 non-null  float64
8   Absent_R                             17388 non-null  float64
9   Neutral masked                       17396 non-null  float64
10  INITIAL CONFIDENCE GRADING           17396 non-null  int32
11  FINAL CONFIDENCE GRADING             17396 non-null  int32
dtypes: float64(8), int32(4)
memory usage: 1.3 MB
```

In [63]: `df.isnull().sum()`

Out[63]:

Drug	0
Tier	48
Common Variant	0
Present_SOLO_R	8
Present_SOLO_SR	8
Present_S	8
Absent_S	8
Present_R	8
Absent_R	8
Neutral masked	0
INITIAL CONFIDENCE GRADING	0
FINAL CONFIDENCE GRADING	0

dtype: int64

In [60]: `df.dtypes`

```
Out[60]: Drug                int32
Tier                float64
Common Variant      int32
Present_SOLO_R      float64
Present_SOLO_SR      float64
Present_S            float64
Absent_S             float64
Present_R            float64
Absent_R             float64
Neutral masked       float64
INITIAL CONFIDENCE GRADING  int32
FINAL CONFIDENCE GRADING  int32
dtype: object
```

```
In [61]: df.describe()
```

Out[61]:

	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_
count	17396.000000	17348.000000	17396.000000	17388.000000	17388.000000	17388.000000	17388.000000	17388.000000	17388.000000
mean	7.132847	0.415610	6947.03863	0.000279	0.000954	0.001940	0.644772	0.000476	0.214600
std	4.349835	0.492841	4133.50626	0.010180	0.016847	0.028009	0.248664	0.011424	0.284070
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	5.000000	0.000000	3183.750000	0.000000	0.000000	0.000118	0.422117	0.000000	0.016800
50%	7.000000	0.000000	7145.000000	0.000000	0.000000	0.000118	0.587725	0.000000	0.079400
75%	12.000000	1.000000	10673.250000	0.000000	0.000138	0.000236	0.861942	0.000000	0.379900
max	14.000000	1.000000	13448.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
In [64]: for feature in df.columns:
median_value = df[feature].median()
df[feature].fillna(median_value, inplace=True)
```

```
In [65]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17396 entries, 0 to 17395
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Drug                  17396 non-null  int32
1   Tier                  17396 non-null  float64
2   Common Variant        17396 non-null  int32
3   Present_SOLO_R        17396 non-null  float64
4   Present_SOLO_SR        17396 non-null  float64
5   Present_S              17396 non-null  float64
6   Absent_S              17396 non-null  float64
7   Present_R              17396 non-null  float64
8   Absent_R              17396 non-null  float64
9   Neutral masked         17396 non-null  float64
10  INITIAL CONFIDENCE GRADING  17396 non-null  int32
11  FINAL CONFIDENCE GRADING  17396 non-null  int32
dtypes: float64(8), int32(4)
memory usage: 1.3 MB
```

```
In [67]: df.corr()
```

Out[67]:

	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R	Neutr masked
Drug	1.000000	-0.093594	0.107236	0.012414	0.012896	-0.000168	0.036932	0.017281	0.231247	0.000000
Tier	-0.093594	1.000000	-0.076048	-0.021700	-0.007533	0.018833	0.100121	-0.031870	-0.569943	0.028000
Common Variant	0.107236	-0.076048	1.000000	0.010508	0.007713	0.007509	-0.179507	0.007627	0.030117	0.007400
Present_SOLO_R	0.012414	-0.021700	0.010508	1.000000	0.631376	0.020307	0.008458	0.871117	0.006031	0.001700
Present_SOLO_SR	0.012896	-0.007533	0.007713	0.631376	1.000000	0.413592	-0.022704	0.668213	-0.006268	0.207200
Present_S	-0.000168	0.018833	0.007509	0.020307	0.413592	1.000000	-0.085283	0.299541	-0.031625	0.480600
Absent_S	0.036932	0.100121	-0.179507	0.008458	-0.022704	-0.085283	1.000000	-0.008724	0.373310	-0.084700
Present_R	0.017281	-0.031870	0.007627	0.871117	0.668213	0.299541	-0.008724	1.000000	0.013122	0.107500
Absent_R	0.231247	-0.569943	0.030117	0.006031	-0.006268	-0.031625	0.373310	0.013122	1.000000	-0.043600
Neutral masked	0.000060	0.028058	0.007443	0.001739	0.207243	0.480660	-0.084763	0.107510	-0.043654	1.000000
INITIAL CONFIDENCE GRADING	-0.100572	0.089321	-0.064167	-0.175689	-0.200729	-0.197813	0.075625	-0.191941	0.005752	-0.396600
FINAL CONFIDENCE GRADING	-0.117773	0.162170	-0.048383	-0.112212	0.040371	0.265929	0.023106	-0.030491	-0.096255	0.549800

In []:

```
from sklearn.preprocessing import MinMaxScaler

# numerical columns for normalizing
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

# Min-Max scaling (scaling features between 0 and 1)
min_max_scaler = MinMaxScaler()
df[numerical_columns] = min_max_scaler.fit_transform(df[numerical_columns])
```

In [69]: df[numerical_columns]

Out[69]:

	Tier	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R	Neutral masked
0	0.0	0.128571	0.133776	0.005898	0.606014	0.099019	0.028618	0.0
1	0.0	0.003922	0.010663	0.006016	0.283827	0.003374	0.051825	0.0
2	0.0	0.000700	0.000969	0.000236	0.607874	0.000633	0.105125	0.0
3	0.0	0.000560	0.001800	0.001180	0.607564	0.000527	0.105207	0.0
4	1.0	0.000420	0.013848	0.011443	0.604192	0.000316	0.011562	0.0
...
17391	0.0	0.000280	0.036837	0.033620	0.240081	0.002109	0.208446	1.0
17392	0.0	0.000000	0.000000	0.042350	0.237213	0.002004	0.208528	1.0
17393	1.0	0.000000	0.000000	0.005898	0.249186	0.000000	0.006970	1.0
17394	1.0	0.000000	0.000000	0.007432	0.248683	0.000000	0.006970	1.0
17395	0.0	0.000000	0.000000	0.000000	0.362252	0.000844	0.379418	0.0

17396 rows × 8 columns

```
# Selecting categorical columns
categorical_cols = df[['Drug', 'Common Variant', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']]

# Selecting numerical columns
numerical_columns = df[['Tier', 'Present_SOLO_R', 'Present_SOLO_SR',
                        'Present_S', 'Absent_S', 'Present_R', 'Absent_R', 'Neutral masked']]

# Concatenate along columns(axis=1) to merge them
df_merged = pd.concat([categorical_cols, numerical_columns], axis=1)

# Print the merged DataFrame
df_merged
```

Out[70]:

	Drug	Common Variant	INITIAL CONFIDENCE GRADING	FINAL CONFIDENCE GRADING	Tier	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R
0	0	11986	0	0	0.0	0.128571	0.133776	0.005898	0.606014	0.099019	0.028618
1	0	2333	0	0	0.0	0.003922	0.010663	0.006016	0.283827	0.003374	0.051825
2	0	12271	0	1	0.0	0.000700	0.000969	0.000236	0.607874	0.000633	0.105125
3	0	12111	4	1	0.0	0.000560	0.001800	0.001180	0.607564	0.000527	0.105207
4	0	12971	4	2	1.0	0.000420	0.013848	0.011443	0.604192	0.000316	0.011562
...
17391	14	13290	2	4	0.0	0.000280	0.036837	0.033620	0.240081	0.002109	0.208446
17392	14	12149	2	4	0.0	0.000000	0.000000	0.042350	0.237213	0.002004	0.208528
17393	14	11603	2	4	1.0	0.000000	0.000000	0.005898	0.249186	0.000000	0.006970
17394	14	12784	2	4	1.0	0.000000	0.000000	0.007432	0.248683	0.000000	0.006970
17395	14	11566	5	5	0.0	0.000000	0.000000	0.000000	0.362252	0.000844	0.379418

17396 rows × 12 columns

In [71]:

```
#from sklearn.preprocessing import OneHotEncoder

# Initialize the OneHotEncoder
#encoder = OneHotEncoder()

# Fit and transform the categorical columns
#categorical_cols = ['Drug', 'Common Variant', 'INITIAL CONFIDENCE GRADING', 'FINAL CONFIDENCE GRADING']
#encoded_data = encoder.fit_transform(df[categorical_cols])

# Convert the encoded data to a DataFrame
#encoded_df = pd.DataFrame(encoded_data.toarray(), columns=encoder.get_feature_names_out(categorical_cols))

# Drop the original categorical columns from the DataFrame
#df.drop(columns=categorical_cols, inplace=True)

# Concatenate the encoded DataFrame with the original DataFrame
#df_encoded = pd.concat([df, encoded_df], axis=1)
```

In [73]:

```
# Swap the values of "final confidence" and "initial confidence" in the last row
last_index = df.index[-1] # Get the index of the last row
final_confidence = df.at[last_index, 'FINAL CONFIDENCE GRADING']
initial_confidence = df.at[last_index, 'INITIAL CONFIDENCE GRADING']

# Swap the values
df.at[last_index, 'INITIAL CONFIDENCE GRADING'] = final_confidence
df.at[last_index, 'FINAL CONFIDENCE GRADING'] = initial_confidence

# Print the DataFrame to verify the changes
print(df)
```


	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S \
0	0	0.0	11986	0.128571	0.133776	0.005898
1	0	0.0	2333	0.003922	0.010663	0.006016
2	0	0.0	12271	0.000700	0.000969	0.000236
3	0	0.0	12111	0.000560	0.001800	0.001180
4	0	1.0	12971	0.000420	0.013848	0.011443
...
17391	14	0.0	13290	0.000280	0.036837	0.033620
17392	14	0.0	12149	0.000000	0.000000	0.042350
17393	14	1.0	11603	0.000000	0.000000	0.005898
17394	14	1.0	12784	0.000000	0.000000	0.007432
17395	14	0.0	11566	0.000000	0.000000	0.000000

	Absent_S	Present_R	Absent_R	Neutral masked \
0	0.606014	0.099019	0.028618	0.0
1	0.283827	0.003374	0.051825	0.0
2	0.607874	0.000633	0.105125	0.0
3	0.607564	0.000527	0.105207	0.0
4	0.604192	0.000316	0.011562	0.0
...
17391	0.240081	0.002109	0.208446	1.0
17392	0.237213	0.002004	0.208528	1.0
17393	0.249186	0.000000	0.006970	1.0
17394	0.248683	0.000000	0.006970	1.0
17395	0.362252	0.000844	0.379418	0.0

	INITIAL CONFIDENCE GRADING	FINAL CONFIDENCE GRADING
0	0	0
1	0	0
2	0	1
3	4	1
4	4	2
...
17391	2	4
17392	2	4
17393	2	4
17394	2	4
17395	5	5

[17396 rows x 12 columns]

In [41]: `df.head()`

Out[41]:

	Drug	Tier	Common Variant	Present_SOLO_R	Present_SOLO_SR	Present_S	Absent_S	Present_R	Absent_R	Neutral masked	INITIAL CONFIDENCE GRADING	FINAL CONFIDENCE GRADING
0	0	0.0	11986	0.128571	0.133776	0.005898	0.606014	0.099019	0.028618	0.0	0	0
1	0	0.0	2333	0.003922	0.010663	0.006016	0.283827	0.003374	0.051825	0.0	0	0
2	0	0.0	12271	0.000700	0.000969	0.000236	0.607874	0.000633	0.105125	0.0	0	0
3	0	0.0	12111	0.000560	0.001800	0.001180	0.607564	0.000527	0.105207	0.0	4	4
4	0	1.0	12971	0.000420	0.013848	0.011443	0.604192	0.000316	0.011562	0.0	4	4

In [74]: `# Check for specific values in the 'Drug' column`
`drug_counts = df['Drug'].value_counts()`
`print("Counts for each drug:")`
`print(drug_counts)`

Counts for each drug:

```

5    2701
7    2430
13   1714
0    1571
14   1522
6    1438
12   1144
2    1067
9     703
8     615
3     608
11    580
1     537
10    406
4     360
Name: Drug, dtype: int64

```

In [75]: `# Check for specific values in the 'Tier' column`
`tier_counts = df['Tier'].value_counts()`
`print("\nCounts for each tier:")`
`print(tier_counts)`

Counts for each tier:

```

0.0    10186
1.0     7210
Name: Tier, dtype: int64

```

```
In [76]: INITIAL_CONFIDENCE_GRADING = df['INITIAL CONFIDENCE GRADING'].value_counts()
```

```
# Print the counts for INITIAL CONFIDENCE GRADING
print("Counts for INITIAL CONFIDENCE GRADING:")
print(INITIAL_CONFIDENCE_GRADING)
```

Counts for INITIAL CONFIDENCE GRADING:

```
4    16805
0      226
2      222
1       88
5       40
3       15
```

Name: INITIAL CONFIDENCE GRADING, dtype: int64

```
In [77]: FINAL_CONFIDENCE_GRADING = df['FINAL CONFIDENCE GRADING'].value_counts()
```

```
# Print the counts for INITIAL CONFIDENCE GRADING
print("Counts for FINAL CONFIDENCE GRADING:")
print(FINAL_CONFIDENCE_GRADING)
```

Counts for FINAL CONFIDENCE GRADING:

```
2    15910
1     1004
4      213
0      196
5       40
3       33
```

Name: FINAL CONFIDENCE GRADING, dtype: int64

```
In [78]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
# Step 1: Data Preprocessing
```

```
X = df.drop(columns=['FINAL CONFIDENCE GRADING'])
y = df['FINAL CONFIDENCE GRADING']
```

```
In [81]: # Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 2: Model Selection and Training
```

```
model = RandomForestClassifier() # You can use any other classifier here
model.fit(X_train, y_train)
```

```
# Step 3: Model Evaluation
```

```
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9804597701149426

```
In [82]: from sklearn.metrics import classification_report
```

```
# Here, y_test and y_pred are the true labels and predicted labels respectively
```

```
# Print classification report
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	46
1	0.84	0.84	0.84	203
2	0.99	0.99	0.99	3180
3	1.00	0.75	0.86	4
4	1.00	1.00	1.00	40
5	1.00	0.71	0.83	7
accuracy			0.98	3480
macro avg	0.96	0.88	0.91	3480
weighted avg	0.98	0.98	0.98	3480

```
In [83]: from sklearn.metrics import accuracy_score
```

```
# Assuming you have already trained your model and obtained predictions (y_pred) on the testing set (X_test)
# Here, model is your trained classifier
```

```
# Make predictions on the testing set
```

```
y_pred = model.predict(X_test)
```

```
# Calculate accuracy on the testing set
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy on testing set:", accuracy)
```

Accuracy on testing set: 0.9804597701149426

```
In [88]: from sklearn.ensemble import RandomForestClassifier
```

```
# X_train and y_train are the training features and labels
```

```

rf_classifier = RandomForestClassifier()

# Fit the model
rf_classifier.fit(X_train, y_train)

# Get feature importances
feature_importances = rf_classifier.feature_importances_

```

```

In [85]: # Match feature importances with feature names
feature_importance_dict = dict(zip(X_train.columns, feature_importances))

# Sort feature importances in descending order
sorted_feature_importances = sorted(feature_importance_dict.items(), key=lambda x: x[1], reverse=True)

# Print feature importances
for feature, importance in sorted_feature_importances:
    print(f"Feature: {feature}, Importance: {importance}")

```

```

Feature: Common Variant, Importance: 0.3616054767513026
Feature: INITIAL CONFIDENCE GRADING, Importance: 0.18772829410611908
Feature: Present_SOLO_R, Importance: 0.09160159071097614
Feature: Absent_R, Importance: 0.07203442353458094
Feature: Present_R, Importance: 0.06035632406647119
Feature: Present_S, Importance: 0.05831126068993332
Feature: Absent_S, Importance: 0.05264770211189792
Feature: Neutral masked, Importance: 0.04520675704198554
Feature: Present_SOLO_SR, Importance: 0.032340385919536425
Feature: Drug, Importance: 0.031238420898142365
Feature: Tier, Importance: 0.006929364169054611

```

```

In [91]: # Select the most important features identified
important_features = ['Common Variant', 'INITIAL CONFIDENCE GRADING',
                     'Present_SOLO_R', 'Neutral masked', 'Present_R', 'Present_S', 'Absent_R', 'Absent_S']

# Analyze the relationship between each important feature and final confidence gradings
for feature in important_features:
    # Group the data by the feature and calculate the mean final confidence grading for each group
    feature_confidence_mean = df.groupby(feature)['FINAL CONFIDENCE GRADING'].mean()

    # Print the relationship between the feature and final confidence gradings
    print(f"\nRelationship between '{feature}' and Final Confidence Gradings:\n")
    print(feature_confidence_mean)

```

Relationship between 'Common Variant' and Final Confidence Gradings:

```

Common Variant
0      2.0
1      2.0
2      2.0
3      2.0
4      2.0
...
13444  2.0
13445  2.0
13446  2.0
13447  2.0
13448  2.0
Name: FINAL CONFIDENCE GRADING, Length: 13449, dtype: float64

```

Relationship between 'INITIAL CONFIDENCE GRADING' and Final Confidence Gradings:

```

INITIAL CONFIDENCE GRADING
0      0.141593
1      1.000000
2      3.959459
3      3.000000
4      1.947694
5      5.000000
Name: FINAL CONFIDENCE GRADING, dtype: float64

```

Relationship between 'Present_SOLO_R' and Final Confidence Gradings:

```

Present_SOLO_R
0.000000    2.003617
0.000140    1.714411
0.000280    1.629771
0.000420    1.657895
0.000560    1.633803
...
0.158683    0.000000
0.216246    0.000000
0.243417    0.000000
0.746779    0.000000
1.000000    0.000000
Name: FINAL CONFIDENCE GRADING, Length: 92, dtype: float64

```

Relationship between 'Neutral masked' and Final Confidence Gradings:

```
Neutral masked
0.0      1.927099
1.0      3.959459
Name: FINAL CONFIDENCE GRADING, dtype: float64
```

Relationship between 'Present_R' and Final Confidence Gratings:

```
Present_R
0.000000    1.995467
0.000105    1.805730
0.000211    1.788214
0.000316    1.864035
0.000422    1.803922
...
0.236739    0.000000
0.241485    0.000000
0.540124    4.000000
0.689233    0.000000
1.000000    0.000000
Name: FINAL CONFIDENCE GRADING, Length: 152, dtype: float64
```

Relationship between 'Present_S' and Final Confidence Gratings:

```
Present_S
0.000000    1.712017
0.000118    1.967196
0.000236    1.970423
0.000354    1.954918
0.000472    1.952278
...
0.908104    4.000000
0.908458    4.000000
0.976997    4.000000
0.984664    4.000000
1.000000    4.000000
Name: FINAL CONFIDENCE GRADING, Length: 253, dtype: float64
```

Relationship between 'Absent_R' and Final Confidence Gratings:

```
Absent_R
0.000000    4.0
0.000082    4.0
0.000246    4.0
0.000574    4.0
0.000738    4.0
...
0.999508    5.0
0.999672    5.0
0.999754    5.0
0.999918    5.0
1.000000    5.0
Name: FINAL CONFIDENCE GRADING, Length: 535, dtype: float64
```

Relationship between 'Absent_S' and Final Confidence Gratings:

```
Absent_S
0.000000    4.000000
0.000116    4.000000
0.000155    4.000000
0.000271    4.000000
0.000426    4.000000
...
0.999845    2.000000
0.999884    1.987421
0.999923    2.000000
0.999961    2.000000
1.000000    2.021898
Name: FINAL CONFIDENCE GRADING, Length: 748, dtype: float64
```

```
In [92]: from scipy import stats

# Select the most important features identified
important_features = ['Common Variant', 'INITIAL CONFIDENCE GRADING',
                     'Present_SOLO_R', 'Neutral masked', 'Present_R', 'Present_S', 'Absent_R', 'Absent_S']

# Analyze the relationship between each important feature and final confidence gradings
for feature in important_features:
    # Find pointbserialr correlation coefficient
    correlation_coefficient, p_value = stats.pointbserialr(df[feature], df['FINAL CONFIDENCE GRADING'])

    # Print the results
    print(f"\nPoint-biserial correlation coefficient between '{feature}' and 'FINAL CONFIDENCE GRADING': {correlation_coefficient}")
    print(f"P-value: {p_value}")
```

Point-biserial correlation coefficient between 'Common Variant' and 'FINAL CONFIDENCE GRADING': -0.04838250627849848
P-value: 1.7180372270874056e-10

Point-biserial correlation coefficient between 'INITIAL CONFIDENCE GRADING' and 'FINAL CONFIDENCE GRADING': 0.2797176016094426
P-value: 4.0578401675859e-310

Point-biserial correlation coefficient between 'Present_SOLO_R' and 'FINAL CONFIDENCE GRADING': -0.11221237922196936
P-value: 7.387710645132828e-50

Point-biserial correlation coefficient between 'Neutral masked' and 'FINAL CONFIDENCE GRADING': 0.5498151025126343
P-value: 0.0

Point-biserial correlation coefficient between 'Present_R' and 'FINAL CONFIDENCE GRADING': -0.03049148049458477
P-value: 5.763592431595842e-05

Point-biserial correlation coefficient between 'Present_S' and 'FINAL CONFIDENCE GRADING': 0.26592890781587236
P-value: 2.162312346552411e-279

Point-biserial correlation coefficient between 'Absent_R' and 'FINAL CONFIDENCE GRADING': -0.0962551796587794
P-value: 4.343903697446709e-37

Point-biserial correlation coefficient between 'Absent_S' and 'FINAL CONFIDENCE GRADING': 0.02310558389647338
P-value: 0.0023062047405744313

In [93]: `import matplotlib.pyplot as plt`

```
# Select the most important features identified
important_features = ['Common Variant', 'INITIAL CONFIDENCE GRADING',
                     'Present_SOLO_R', 'Neutral masked', 'Present_R', 'Present_S', 'Absent_R', 'Absent_S']

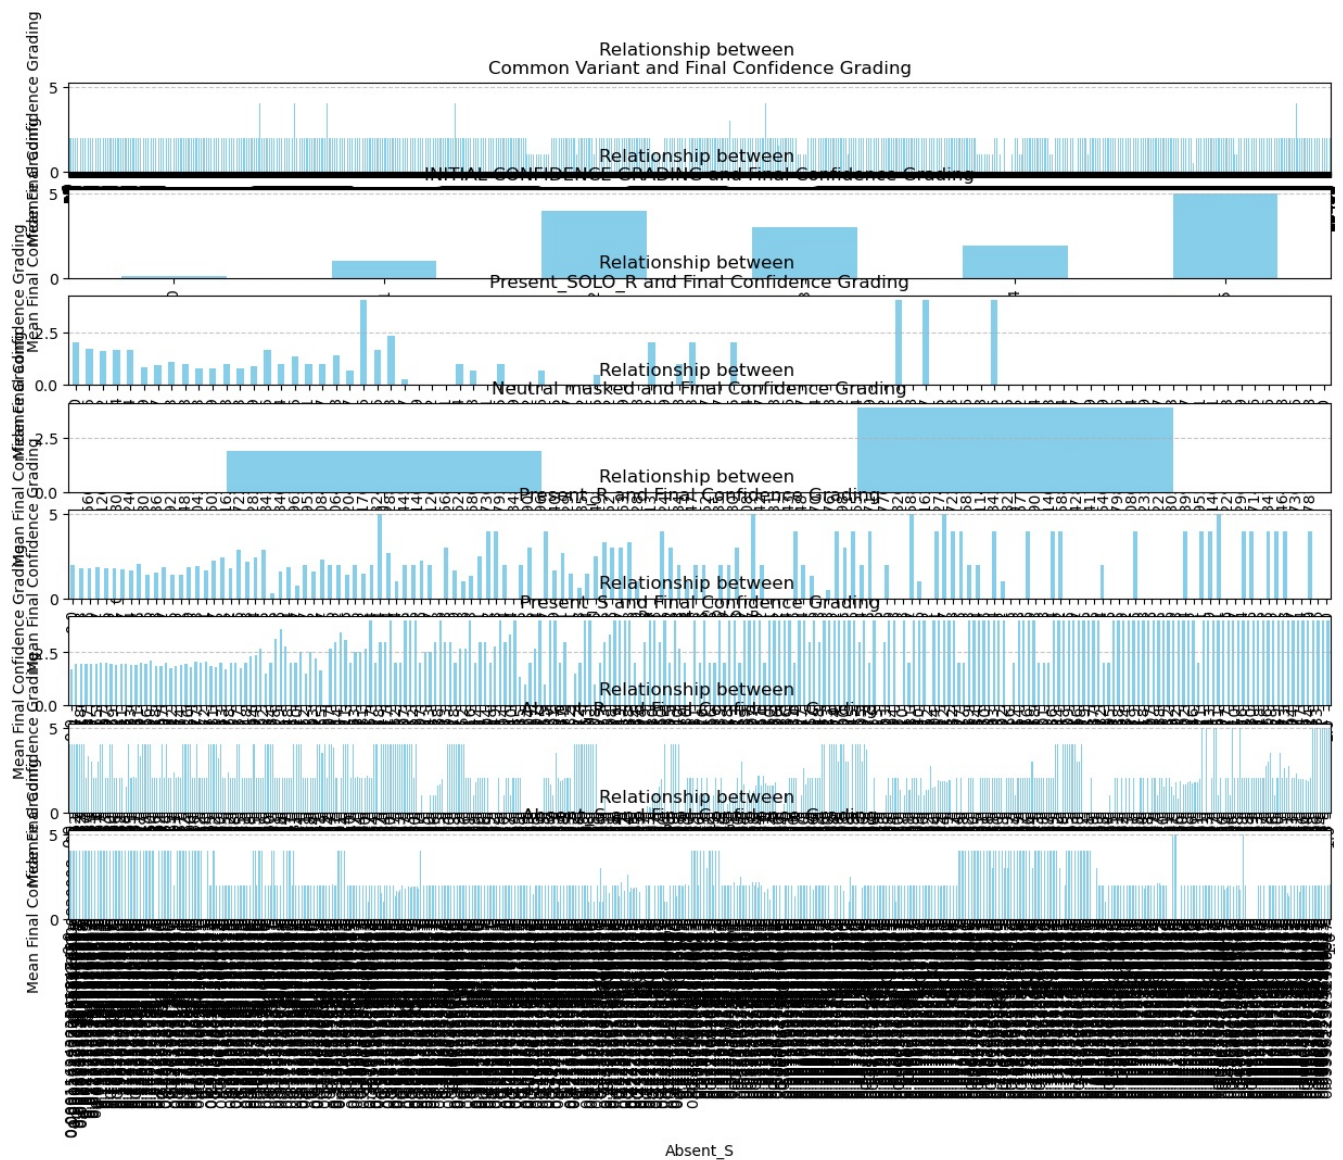
# Set up subplots for each feature
fig, axes = plt.subplots(nrows=len(important_features), figsize=(15, 10))

# Analyze the relationship between each important feature and final confidence gradings
for i, feature in enumerate(important_features):
    # Group the data by the feature and calculate the mean final confidence grading for each group
    feature_confidence_mean = df.groupby(feature)['FINAL CONFIDENCE GRADING'].mean()

    # Plot the bar plot
    feature_confidence_mean.plot(kind='bar', ax=axes[i], color='skyblue')
    axes[i].set_title(f'Relationship between \n{feature} and Final Confidence Grading')
    axes[i].set_ylabel('Mean Final Confidence Grading')
    axes[i].set_xlabel(feature)
    axes[i].grid(axis='y', linestyle='--', alpha=0.95)

plt.tight_layout()
plt.show()
```

C:\Users\hp\AppData\Local\Temp\ipykernel_12104\3065845851.py:22: UserWarning: Tight layout not applied. tight_layout cannot make axes height small enough to accommodate all axes decorations.
plt.tight_layout()



In []:

In []: