

```
In [30]: import pandas as pd
url = ("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-S
df = pd.read_csv(url,header = 0)
```

```
In [31]: df.head()
```

```
Out[31]:
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	horsepower	peak-rpm	
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111.0	5000.0	
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111.0	5000.0	
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154.0	5000.0	
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102.0	5500.0	
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115.0	5500.0	

5 rows × 29 columns

```
In [32]: df.columns
```

```
Out[32]: Index(['symboling', 'normalized-losses', 'make', 'aspiration', 'num-of-doors',
              'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
              'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',
              'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',
              'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price',
              'city-L/100km', 'horsepower-binned', 'diesel', 'gas'],
              dtype='object')
```

```
In [33]: cols = list(df.columns)

# move the outcome column to the end
cols.append(cols.pop(cols.index('price')))

# reorder the columns in the dataframe
df = df[cols]

# save the updated dataframe back to the CSV file
df.to_csv('url', index=False)
```

```
In [ ]:
```

```
In [29]: df.describe()
```

```
Out[29]:
```

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size	bore	stroke	compression-ratio
count	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	197.000000	201.000000
mean	0.840796	122.000000	98.797015	0.837102	0.915126	53.766667	2555.666667	126.875622	3.330692	3.256904	10.164000
std	1.254802	31.99625	6.066366	0.059213	0.029187	2.447822	517.296727	41.546834	0.268072	0.319256	4.004000
min	-2.000000	65.000000	86.600000	0.678039	0.837500	47.800000	1488.000000	61.000000	2.540000	2.070000	7.000000
25%	0.000000	101.000000	94.500000	0.801538	0.890278	52.000000	2169.000000	98.000000	3.150000	3.110000	8.600000
50%	1.000000	122.000000	97.000000	0.832292	0.909722	54.100000	2414.000000	120.000000	3.310000	3.290000	9.000000
75%	2.000000	137.000000	102.400000	0.881788	0.925000	55.500000	2926.000000	141.000000	3.580000	3.410000	9.400000
max	3.000000	256.000000	120.900000	1.000000	1.000000	59.800000	4066.000000	326.000000	3.940000	4.170000	23.000000

```
In [23]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: import pandas as pd

# Select the object features to convert
object_features = ['feature1', 'feature2', 'feature3']

# Convert the selected object features to float
df[object_features] = df[object_features].astype(float)

# Save the updated dataframe to a new CSV file
df.to_csv('updated_data.csv', index=False)
```

```
In [4]: df.dtypes
```

```
Out[4]: symboling          int64
normalized-losses      int64
make                   object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                 float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders        object
engine-size            int64
fuel-system            object
bore                  float64
stroke                float64
compression-ratio      float64
horsepower             float64
peak-rpm              float64
city-mpg              int64
highway-mpg           int64
price                 float64
city-L/100km          float64
horsepower-binned      object
diesel                int64
gas                   int64
dtype: object
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 201 entries, 0 to 200
Data columns (total 29 columns):
#   Column              Non-Null Count  Dtype
---  -
0   symboling            201 non-null    int64
1   normalized-losses    201 non-null    int64
2   make                 201 non-null    object
3   aspiration           201 non-null    object
4   num-of-doors         201 non-null    object
5   body-style           201 non-null    object
6   drive-wheels         201 non-null    object
7   engine-location      201 non-null    object
8   wheel-base          201 non-null    float64
9   length              201 non-null    float64
10  width               201 non-null    float64
11  height              201 non-null    float64
12  curb-weight          201 non-null    int64
13  engine-type          201 non-null    object
14  num-of-cylinders     201 non-null    object
15  engine-size          201 non-null    int64
16  fuel-system          201 non-null    object
17  bore                 201 non-null    float64
18  stroke              197 non-null    float64
19  compression-ratio    201 non-null    float64
20  horsepower           201 non-null    float64
21  peak-rpm            201 non-null    float64
22  city-mpg            201 non-null    int64
23  highway-mpg         201 non-null    int64
24  price               201 non-null    float64
25  city-L/100km        201 non-null    float64
26  horsepower-binned    200 non-null    object
27  diesel              201 non-null    int64
28  gas                 201 non-null    int64
dtypes: float64(11), int64(8), object(10)
memory usage: 45.7+ KB
```

```
In [ ]:
```

```
In [ ]:
```

```
In [6]: df.corr()
```

Out[6]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size	bore	stroke	compression-ratio	horsepower
symboling	1.000000	0.466264	-0.535987	-0.365404	-0.242423	-0.550160	-0.233118	-0.110581	-0.140019	-0.008245	-0.182196	
normalized-losses	0.466264	1.000000	-0.056661	0.019424	0.086802	-0.373737	0.099404	0.112360	-0.029862	0.055563	-0.114713	
wheel-base	-0.535987	-0.056661	1.000000	0.876024	0.814507	0.590742	0.782097	0.572027	0.493244	0.158502	0.250313	
length	-0.365404	0.019424	0.876024	1.000000	0.857170	0.492063	0.880665	0.685025	0.608971	0.124139	0.159733	
width	-0.242423	0.086802	0.814507	0.857170	1.000000	0.306002	0.866201	0.729436	0.544885	0.188829	0.189867	
height	-0.550160	-0.373737	0.590742	0.492063	0.306002	1.000000	0.307581	0.074694	0.180449	-0.062704	0.259737	
curb-weight	-0.233118	0.099404	0.782097	0.880665	0.866201	0.307581	1.000000	0.849072	0.644060	0.167562	0.156433	
engine-size	-0.110581	0.112360	0.572027	0.685025	0.729436	0.074694	0.849072	1.000000	0.572609	0.209523	0.028889	
bore	-0.140019	-0.029862	0.493244	0.608971	0.544885	0.180449	0.644060	0.572609	1.000000	-0.055390	0.001263	
stroke	-0.008245	0.055563	0.158502	0.124139	0.188829	-0.062704	0.167562	0.209523	-0.055390	1.000000	0.187923	
compression-ratio	-0.182196	-0.114713	0.250313	0.159733	0.189867	0.259737	0.156433	0.028889	0.001263	0.187923	1.000000	
horsepower	0.075819	0.217299	0.371147	0.579821	0.615077	-0.087027	0.757976	0.822676	0.566936	0.098462	-0.214514	1.000000
peak-rpm	0.279740	0.239543	-0.360305	-0.285970	-0.245800	-0.309974	-0.279361	-0.256733	-0.267392	-0.065713	-0.435780	
city-mpg	-0.035527	-0.225016	-0.470606	-0.665192	-0.633531	-0.049800	-0.749543	-0.650546	-0.582027	-0.034696	0.331425	
highway-mpg	0.036233	-0.181877	-0.543304	-0.698142	-0.680635	-0.104812	-0.794889	-0.679571	-0.591309	-0.035201	0.268465	
price	-0.082391	0.133999	0.584642	0.690628	0.751265	0.135486	0.834415	0.872335	0.543155	0.082310	0.071107	
city-L/100km	0.066171	0.238567	0.476153	0.657373	0.673363	0.003811	0.785353	0.745059	0.554610	0.037300	-0.299372	
diesel	-0.196735	-0.101546	0.307237	0.211187	0.244356	0.281578	0.221046	0.070779	0.054458	0.241303	0.985231	
gas	0.196735	0.101546	-0.307237	-0.211187	-0.244356	-0.281578	-0.221046	-0.070779	-0.054458	-0.241303	-0.985231	

In [7]: #Find the correlation between the following columns: bore, stroke, compression-ratio, and horsepower

```
df[['bore', 'stroke', 'compression-ratio', 'horsepower']].corr()
```


Out[7]:

	bore	stroke	compression-ratio	horsepower
bore	1.000000	-0.055390	0.001263	0.566936
stroke	-0.055390	1.000000	0.187923	0.098462
compression-ratio	0.001263	0.187923	1.000000	-0.214514
horsepower	0.566936	0.098462	-0.214514	1.000000

In [8]: # Engine size as potential predictor variable of price

```
sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

Out[8]: (0.0, 53042.53513934235)




In [9]: df[["engine-size", "price"]].corr()

```
Out[9]:
```

	engine-size	price
engine-size	1.000000	0.872335
price	0.872335	1.000000


```
In [10]: sns.regplot(x="highway-mpg", y="price", data=df)
```

```
Out[10]: <AxesSubplot:xlabel='highway-mpg', ylabel='price'>
```



```
In [11]: sns.regplot(x="peak-rpm", y="price", data=df)
```

```
Out[11]: <AxesSubplot:xlabel='peak-rpm', ylabel='price'>
```



```
In [12]: df[['peak-rpm', 'price']].corr()
```

```
Out[12]:
```

	peak-rpm	price
peak-rpm	1.000000	-0.101616
price	-0.101616	1.000000

```
In [13]: df[['stroke', 'price']].corr()
```


```
Out[13]:
```

	stroke	price
stroke	1.00000	0.08231
price	0.08231	1.00000

```
In [14]: sns.regplot(x="stroke", y="price", data=df)
```


```
plt.ylim(0,)
```

```
Out[14]: (0.0, 47414.1)
```




```
In [15]: sns.boxplot(x="body-style", y="price", data=df)
```

```
Out[15]: <AxesSubplot:xlabel='body-style', ylabel='price'>
```



```
In [16]: sns.boxplot(x="drive-wheels", y="price", data=df)
```

```
Out[16]: <AxesSubplot:xlabel='drive-wheels', ylabel='price'>
```



In [17]: `df.describe()`

Out[17]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size	bore	stroke	compression-ratio
count	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	197.000000	201.000000
mean	0.840796	122.000000	98.797015	0.837102	0.915126	53.766667	2555.666667	126.875622	3.330692	3.256904	10.164383
std	1.254802	31.99625	6.066366	0.059213	0.029187	2.447822	517.296727	41.546834	0.268072	0.319256	4.004515
min	-2.000000	65.000000	86.600000	0.678039	0.837500	47.800000	1488.000000	61.000000	2.540000	2.070000	7.000000
25%	0.000000	101.000000	94.500000	0.801538	0.890278	52.000000	2169.000000	98.000000	3.150000	3.110000	8.600000
50%	1.000000	122.000000	97.000000	0.832292	0.909722	54.100000	2414.000000	120.000000	3.310000	3.290000	9.000000
75%	2.000000	137.000000	102.400000	0.881788	0.925000	55.500000	2926.000000	141.000000	3.580000	3.410000	9.400000
max	3.000000	256.000000	120.900000	1.000000	1.000000	59.800000	4066.000000	326.000000	3.940000	4.170000	23.000000

In [18]: `df.describe(include=['object'])`

Out[18]:

	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	engine-type	num-of-cylinders	fuel-system	horsepower-binned
count	201	201	201	201	201	201	201	201	201	200
unique	22	2	2	5	3	2	6	7	8	3
top	toyota	std	four	sedan	fwd	front	ohc	four	mpfi	Low
freq	32	165	115	94	118	198	145	157	92	115

In [19]: `df['drive-wheels'].value_counts()`

Out[19]:

```
fwd    118
rwd     75
4wd      8
Name: drive-wheels, dtype: int64
```

```

In [20]: df['drive-wheels'].value_counts().to_frame()

Out[20]:
      drive-wheels
fwd             118
rwd              75
4wd               8

In [21]: from sklearn.linear_model import LinearRegression

lm = LinearRegression()
lm

Out[21]: LinearRegression()

In [22]: X = df[["length"]]
Y = df["price"]

lm.fit(X,Y)

Out[22]: LinearRegression()

In [23]: Yhat = lm.predict(X)
Yhat[0:10]

Out[23]: array([10801.45044105, 10801.45044105, 11870.44409178, 14275.67980594,
14275.67980594, 14587.46962074, 21446.8455463 , 21446.8455463 ,
21446.8455463 , 14364.76261017])

In [24]: lm.intercept_

Out[24]: -64384.436327421616

In [25]: lm.coef_

Out[25]: array([92690.65779928])

In [75]: Price = -64384.436327421616+92690.65779928*df[['length']]

In [27]: lm1 = LinearRegression()
lm1

LinearRegression()

Out[27]: LinearRegression()

In [28]: lm1.fit(df[["engine-size"]], df[["price"]])
lm1

Out[28]: LinearRegression()

In [29]: lm1.coef_

Out[29]: array([[166.86001569]])

In [30]: lm1.intercept_

Out[30]: array([-7963.33890628])

In [36]: Price = 166.86 * -7963*df[["engine-size"]]

In [37]: Z= df[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']]

In [38]: lm.fit(Z,df["price"])

Out[38]: LinearRegression()

In [39]: lm.intercept_

Out[39]: -15806.624626329198

In [40]: lm.coef_

Out[40]: array([53.49574423,  4.70770099, 81.53026382, 36.05748882])

In [41]: #Price = -15806.624 + 53.49574423horsepower + 4.70770099curb-weight + 81.53026382engine-size + 36.05748882highw

In [42]: lm2 = LinearRegression()
lm2.fit(df[['normalized-losses' , 'highway-mpg']],df['price'])

```

```
Out[42]: LinearRegression()
```

```
In [43]: lm2.coef_
```

```
Out[43]: array([ 1.49789586, -820.45434016])
```

```
In [44]: lm2.intercept_
```

```
Out[44]: 38201.31327245728
```


```
In [45]: Price = 38201.31327245728 +1.49789586normalized-losses-820.45434016highway-mpg
```

```
File "C:\Users\hp\AppData\Local\Temp\ipykernel_11964\2134262372.py", line 1
Price = 38201.31327245728 +1.49789586normalized-losses-820.45434016highway-mpg
^
```

SyntaxError: invalid syntax


```
In [46]: import seaborn as sns
width = 12
height = 10
plt.figure(figsize = (width,height))
sns.regplot(x = "highway-mpg",y = "price", data = df)
plt.ylim(0,)
```

```
Out[46]: (0.0, 48179.76254365848)
```



```
In [47]: plt.figure(figsize=(width, height))
sns.regplot(x="peak-rpm", y="price", data=df)
plt.ylim(0,)
```

```
Out[47]: (0.0, 47414.1)
```

```
In [48]: Y_hat = lm.predict(Z)
```

```
In [49]: plt.figure(figsize=(width, height))
```


```
ax1 = sns.distplot(df['price'], hist=False, color="r", label="Actual Value")
sns.distplot(Y_hat, hist=False, color="b", label="Fitted Values" , ax=ax1)

plt.title('Actual vs Fitted Values for Price')
plt.xlabel('Price (in dollars)')
plt.ylabel('Proportion of Cars')

plt.show()
plt.close()
```

C:\Users\hp\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).
 warnings.warn(msg, FutureWarning)

C:\Users\hp\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).
 warnings.warn(msg, FutureWarning)



```
In [50]: def PlotPolly(model,independent_variable,dependent_variabble,Name):
x_new = np.linspace(15,55,100)
y_new = model(x_new)

plt.plot(independent_variable, dependent_variabble, '.', x_new, y_new, '-')
plt.title ("Polynomial Fit with Matplotlib for Price ~ Length")
ax = plt.gca()
ax.set_facecolor((0.898, 0.898, 0.898))
fig = plt.gcf()
plt.xlabel(Name)
plt.ylabel('Price of Cars')


plt.show()
plt.close()
```

```
In [51]: x = df['highway-mpg']
y = df['price']
```

```
In [54]: import numpy as np
f = np.polyfit(x, y, 3)
p = np.poly1d(f)
print(p)

-1.557 x3 + 204.8 x2 - 8965 x + 1.379e+05
```

```
In [55]: PlotPolly(p, x, y, 'highway-mpg')
```




```
In [56]: np.polyfit(x, y, 3)
```

```
Out[56]: array([-1.55663829e+00,  2.04754306e+02, -8.96543312e+03,  1.37923594e+05])
```

```
In [57]: from sklearn.preprocessing import PolynomialFeatures
```

```
f1 = np.polyfit(x, y, 11)
p1 = np.poly1d(f1)
print(p1)
PlotPolly(p1,x,y, 'Highway MPG')
```

```
-1.243e-08 x11 + 4.722e-06 x10 - 0.0008028 x9 + 0.08056 x8 - 5.297 x7
+ 239.5 x6 - 7588 x5 + 1.684e+05 x4 - 2.565e+06 x3 + 2.551e+07 x2 - 1.491e+08 x + 3.879e+08
```



```
In [58]: pr=PolynomialFeatures(degree=2)
pr
```

```
Out[58]: PolynomialFeatures()
```

```
In [59]: Z_pr=pr.fit_transform(Z)
```

```
In [60]: Z.shape
```

```
Out[60]: (201, 4)
```

```
In [61]: Z_pr.shape
```

Out[61]: (201, 15)

```
In [62]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```
In [63]: Input=[('scale',StandardScaler()), ('polynomial', PolynomialFeatures(include_bias=False)), ('model',LinearRegre
```

```
In [64]: pipe=Pipeline(Input)
pipe
```

```
Out[64]: Pipeline(steps=[('scale', StandardScaler()),
                          ('polynomial', PolynomialFeatures(include_bias=False)),
                          ('model', LinearRegression())])
```


```
In [66]: df[['bore','stroke','compression-ratio','horsepower']].corr()
```

```
Out[66]:
```

	bore	stroke	compression-ratio	horsepower
bore	1.000000	-0.055390	0.001263	0.566936
stroke	-0.055390	1.000000	0.187923	0.098462
compression-ratio	0.001263	0.187923	1.000000	-0.214514
horsepower	0.566936	0.098462	-0.214514	1.000000


```
In [67]: sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

```
Out[67]: (0.0, 53509.36151967784)
```



```
In [68]: sns.boxplot(x="body-style", y="price", data=df)
```

```
Out[68]: <AxesSubplot:xlabel='body-style', ylabel='price'>
```



```
In [69]: df['drive-wheels'].value_counts().to_frame()
```

```
Out[69]:
```

drive-wheels	
fwd	118
rwd	75
4wd	8

```
In [70]: drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'}, inplace=True)
drive_wheels_counts
```

```
Out[70]:
```

value_counts	
fwd	118
rwd	75
4wd	8

```
In [72]: engine_loc_counts = df['engine-location'].value_counts().to_frame()
engine_loc_counts.rename(columns={'engine-location': 'value_counts'}, inplace=True)
engine_loc_counts.index.name = 'engine-location'
engine_loc_counts.head()
```

```
Out[72]:
```

value_counts	
engine-location	
front	198
rear	3

```
In [73]: df_group_one = df[['drive-wheels', 'body-style', 'price']]
```

```
In [74]: # grouping results
df_group_one = df_group_one.groupby(['drive-wheels'],as_index=False).mean()
df_group_one
```

Out[74]:

	drive-wheels	price
0	4wd	10241.000000
1	fwd	9244.779661
2	rwd	19757.613333

```
In [ ]: Use the "groupby" function to find the average "price" of each car based on "body-style".
```

```
In [76]: df_group_one = df[['drive-wheels', 'body-style', 'price']]
```

```
In [77]: df_group_one = df_group_one.groupby(['drive-wheels'],as_index=False).mean()
df_group_one
```

Out[77]:

	drive-wheels	price
0	4wd	10241.000000
1	fwd	9244.779661
2	rwd	19757.613333

```
In [78]: df_gptest = df[['drive-wheels', 'body-style', 'price']]
grouped_test1 = df_gptest.groupby(['drive-wheels', 'body-style'],as_index=False).mean()
grouped_test1
```

Out[78]:

	drive-wheels	body-style	price
0	4wd	hatchback	7603.000000
1	4wd	sedan	12647.333333
2	4wd	wagon	9095.750000
3	fwd	convertible	11595.000000
4	fwd	hardtop	8249.000000
5	fwd	hatchback	8396.387755
6	fwd	sedan	9811.800000
7	fwd	wagon	9997.333333
8	rwd	convertible	23949.600000
9	rwd	hardtop	24202.714286
10	rwd	hatchback	14337.777778
11	rwd	sedan	21711.833333
12	rwd	wagon	16994.222222

```
In [79]: grouped_pivot = grouped_test1.pivot(index='drive-wheels',columns='body-style')
grouped_pivot
```

Out[79]:

	body-style	convertible	hardtop	hatchback	sedan	price
drive-wheels						wagon
	4wd	NaN	NaN	7603.000000	12647.333333	9095.750000
	fwd	11595.0	8249.000000	8396.387755	9811.800000	9997.333333
	rwd	23949.6	24202.714286	14337.777778	21711.833333	16994.222222


```
In [81]: grouped_pivot = grouped_pivot.fillna(0) #fill missing values with 0
grouped_pivot
```

```
Out[81]:
```

	body-style	convertible	hardtop	hatchback	sedan	price
drive-wheels						wagon
	4wd	0.0	0.000000	7603.000000	12647.333333	9095.750000
	fwd	11595.0	8249.000000	8396.387755	9811.800000	9997.333333
	rwd	23949.6	24202.714286	14337.777778	21711.833333	16994.222222

```
In [82]: plt.pcolor(grouped_pivot, cmap='RdBu')
```

```
plt.colorbar()
plt.show()
```



```
In [83]: fig, ax = plt.subplots()
im = ax.pcolor(grouped_pivot, cmap='RdBu')


#label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index

#move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

#rotate label if too long
plt.xticks(rotation=90)

fig.colorbar(im)
plt.show()
```



```
In [84]: from scipy import stats
```

```
In [85]: pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.5846418222655081 with a P-value of P = 8.076488270732989e-20

```
In [86]: pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.809574567003656 with a P-value of P = 6.369057428259557e-48

```
In [87]: pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value )
```

The Pearson Correlation Coefficient is 0.7512653440522674 with a P-value of P = 9.200335510481516e-38

```
In [88]: grouped_test2=df_gptest[['drive-wheels', 'price']].groupby(['drive-wheels'])
grouped_test2.head(2)
```

Out[88]:

	drive-wheels	price
0	rwd	13495.0
1	rwd	16500.0
3	fwd	13950.0
4	4wd	17450.0
5	fwd	15250.0
136	4wd	7603.0

```
In [91]: df_gptest
```

Out[91]:

	drive-wheels	body-style	price
0	rwd	convertible	13495.0
1	rwd	convertible	16500.0
2	rwd	hatchback	16500.0
3	fwd	sedan	13950.0
4	4wd	sedan	17450.0
...
196	rwd	sedan	16845.0
197	rwd	sedan	19045.0
198	rwd	sedan	21485.0
199	rwd	sedan	22470.0
200	rwd	sedan	22625.0

201 rows × 3 columns

```
In [92]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'], grouped_test2.get_group('rwd')['price'],
print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 67.95406500780399 , P = 3.3945443577151245e-23

```
In [93]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'], grouped_test2.get_group('rwd')['price'])
print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 8.580681368924756 , P = 0.004411492211225333

```
In [94]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'], grouped_test2.get_group('fwd')['price'])
print("ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 0.665465750252303 , P = 0.41620116697845666

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```


In []:

In []:

In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js